

Exploring and Visualizing School Achievement and School Effects

Daniel Anderson¹ & Joseph Stevens¹

¹ University of Oregon

Author Note

The research reported here was supported by the Institute of Education Sciences, U.S. Department of Education, through Grant R324C110004, awarded to the University of Oregon. The opinions expressed are those of the authors and do not represent views of the Institute or the U.S. Department of Education.

Correspondence concerning this article should be addressed to Daniel Anderson, 5262 University of Oregon. E-mail: daniela@uoregon.edu

Abstract

Recent years have seen a push for more open data as a means of facilitating transparency and opening doors for further research. Since the implementation of the No Child Left Behind Act, schools and districts across the country have been required to publicly report the number of students scoring in each performance level classification on the statewide test, disaggregated by student subgroups (e.g., gender, race/ethnicity). These data represent a “coarsening” of the underlying continuous data by binning the distributions into categories. The purpose of this paper is to evaluate the extent to which achievement gap estimates from student-level data align with those estimates from coarsened data, using the approach suggested by Ho and Reardon (2012). Following the evaluation, we use this approach to visualize differences in school-level achievement gaps using publicly-available data, and visualize these estimates through geo-spatial mapping. We overlay demographic data as a means to visually evaluate the extent to which demographics of the surrounding area correspond to school-level achievement gaps.

Keywords: Effect Size, Achievement Gaps, Visualization, Maps

Word count: X

Exploring and Visualizing School Achievement and School Effects

Recent years have seen increased attention on open data (i.e., free and publicly available) as one component of open science (e.g., see Dawes, Vidiyasova, & Parkhimovich, 2016; Sieber & Johnson, 2015; Zuiderwijk & Janssen, 2014). The National Institute of Health (NIH), for example, recently made 64 data sharing repositories publicly available (Kittrie & Chattopadhyay, 2016). The National Science Foundation (NSF) has placed a similarly high emphasis on making data from funded projects publicly available, as part of their Open Government Plan (National Science Foundation, 2016). The Institute of Education Sciences (IES) has similarly moved towards greater requirements for data sharing, with Goal 4 Effectiveness grants requiring a data sharing component since 2013, and Goal 3 Efficacy and Replication awards requiring data sharing since 2016 (Institute of Education Sciences, 2016). Despite these movements, data that are fully open remain rare, particularly in educational research. For example, the website data.gov represents the “home of the U.S. Government’s open data”, housing nearly a quarter of a million datasets. Of these datasets, only 365 were tagged with the keyword “Education” (as of this writing), and the vast majority of these receive an openness score of one star out of a possible five.

Privacy issues abound in educational research, given that the research population generally includes minors. Privacy concerns may therefore explain part of the lack of data sharing (although it should be noted that health care and medical research often face similar privacy concerns). One method of avoiding potential issues of confidentiality is to first aggregate the data in some capacity, before releasing the data publicly. This was the approach taken by the *No Child Left Behind Act* (No Child Left Behind, 2002), with data reported at the overall school and district levels, but not at the individual student level. NCLB required annual testing in reading and mathematics in each of grades 3-8, and once in high school, with the results reported at the aggregate levels by student subgroups (e.g., gender, race/ethnicity). Cut scores were required to be established, delineating students into various performance level classifications (PLCs). The state of California, for example,

included the following four categories in order of increasing competency: “Not Met”, “Nearly Met”, “Met”, and “Exceeds”. The proportion of students scoring in each of these categories, by subgroup, were also required to be reported at each of the school and district levels by NCLB, provided a sufficient number of students were represented within each subgroup such that individual students could not be identified.

Evaluating Achievement Gaps With Open Data

Aggregating the data to higher levels, combined with reporting test score data by PLC proportions, helps avoid privacy concerns and allows the data to be reported publicly. Despite mandates from NCLB on public reporting, however, the openness of these data vary by state, in terms of accessing the data for research purposes. Some (e.g., Oregon, California) make the process easy, with downloadable comma separated value (CSV) including the proportion of students scoring in each PLC for all schools and/or districts in the state, by student subgroups. Other states make it relatively easy to access the data for an individual school or district (as mandated by NCLB) but rather difficult to access across schools or districts. Further, the utility of these data for research purposes is not particularly straightforward. While raw percentages scoring within specific PLCs can be compared between schools or between student subgroups, these differences are highly dependent upon the placement of the cut scores (Ho, 2008; Holland, 2002).

Ideally, a transformation could be applied to the “coarsened” data (continuous test score distributions binned into a set of ordinal categories; Reardon & Ho, 2015) to recover the parameters of the full distribution for each subgroup, making the placement of the cut scores inconsequential to evaluating subgroup differences. Ho and Reardon (2012) and Reardon and Ho (2015) propose such an approach, in which the coarsened data are used to construct empirical cumulative distribution functions (ECDFs) for each group. Paired ECDFs for any two groups are then evaluated in the form of a probability-probability (PP) plot, and the area under the PP curve (AUC) provides an estimate of the probability that a randomly selected students from the reference distribution (plotted along the x-axis) would

score higher than a randomly selected student from the focal distribution (plotted along the y-axis; see Ho & Reardon, 2012). This proportion can then be transformed into standard-deviation units, using the following formula, as outlined by Ho (2009)

$$V = \sqrt{2}\Phi^{-1}(AUC) \quad (1)$$

Where Φ^{-1} represents the inverse normal distribution. Under the assumption of respective normality, V is equivalent to Cohen's D , and the dependence upon cut scores is removed, given that it is an estimate of difference between the paired distributions, rather than the difference between proportions at any particular point on the distribution. Importantly, because V can be estimated from coarsened data, it can be used to estimate the overall achievement gap between student subgroups using NCLB mandatorily reported data.

Achievement Gaps in Context

A wealth of previous research has examined achievement differences between student subgroups. Sirin (2005), for example, conducted a meta-analytic review of the relation between socioeconomic status (SES) and academic achievement. The author found that, while there was a medium to strong overall relation, the magnitude was dependent upon (among other factors) the physical location of the school. This is perhaps unsurprising, given that schools draw students from the surrounding neighborhoods, and any disparities between neighborhoods would naturally flow into disparities between schools. When evaluating achievement gaps, it makes sense to consider characteristics of both the school and the surrounding neighborhood/community. For example, it is possible that the racial/ethnic makeup of the surrounding area would relate to achievement gaps at the school, given that identifying as Black in a school with students predominately identifying as White is very different from identifying as Black in a school where students predominately identify as Black (Hanushek, Kain, & Rivkin, 2009).

In addition to school compositional effects, there is considerable evidence that crime rates in the surrounding area relate to student achievement (e.g., Gonzales, Cauce, Friedman,

& Mason, 1996; McCoy, Roy, & Sirkman, 2013). In particular, N. K. Bowen and Bowen (1999) found that measures of neighborhood danger were predictive of a number of outcomes, but particularly measures of behavior and attendance. The authors found that males and students identifying as African American reported higher rates of exposure to both neighborhood and school danger. These results are particularly relevant for achievement gaps when considering Lee and Madyun (2009) found that “Black students were 2.88 times more likely than White peers to reside in the neighborhoods with high crime-high poverty” (p. 165). Further, as noted by Gregory, Skiba, and Noguera (2010), large racial disproportionality exists in terms of school-wide suspension and expulsion rates. This disproportionality may be partially attributable to neighborhood/community differences, which ultimately lead to fewer opportunities for students to learn (i.e., more time spent out of the classroom) and a broadening of the achievement gap.

Summary

The purpose of this paper is to evaluate differences in achievement gaps between schools using open, publicly available data, reported by schools on a mandatory basis as part of NCLB. We begin, however, by first evaluating V using empirical, student-level data collected across the state of Oregon. Specifically, we evaluate the extent to which V corresponds with D when using the full sample. We expect a strong but imperfect relation, given the differing assumption of the two effect sizes. We then manually coarsen the data into the proportion of students scoring in each PLC, re-estimate V , and compare the estimates to those made from the continuous data. Following this investigation, we illustrate how this approach can be useful when working with publicly available data, specifically by using V to calculate achievement gap effect sizes for schools in California and Oregon, combining these data with census data, and producing geo-spatial maps to evaluate both the clustering of school-level estimates of achievement gaps, and the extent to which any clustering relates to the demographics of the surrounding area (e.g., median housing cost).

Method

Data Sources

Multiple sources of data were used. Student-level data included all students in the state of Oregon who took the reading/language arts or mathematics assessments during the 2012-2013 school year. These data are summarized by content area and gender in Table 1, and were collected as part of the National Center on Assessment and Accountability for Special Education (NCAASE; see <http://ncaase.com>), a multi-state collaborative focused on growth modeling and evaluating school effect policies. These data were used to evaluate the correspondence between V and D estimated from the full data, as well as the extent to which V estimated with the discrete (coarsened) data (V_d) corresponded with V estimated from the full, continuous data (V_c)

Publicly-available datasets were used to visualize differences in achievement gaps in both Oregon and California. For both states, data on the proportion of students in each subgroup scoring in each PLC were obtained from the corresponding statewide websites (California Assessment of Student Performance and Progress, 2017; Oregon Department of Education, 2017a). Both states had minimum reporting requirements, although the threshold varied considerably. In California, the minimum reported sample size across student groups was 100, while in Oregon the minimum sample size was only 6. Values below these thresholds in each respective data file were missing.

To map the schools, information on the physical location of the schools, in terms of latitude and longitude, were also necessary. For California, a data file containing this information for all schools in the state was located on the state website (California Department of Education, 2017). For Oregon, no specific information on the latitude and longitude of schools could be located. However, a file containing the school name and physical address was located (Oregon Department of Education, 2017b). These addresses were then transformed to latitude and longitude using Google's geocode application programming interface (API; Google, 2018)

Analyses

All analyses were conducted within the R statistical computing framework (R Core Team, 2017). The *tidyverse* suite of packages (Wickham, 2017) were used for all data preparation and visualization. Effect sizes were estimated using the *esvis* package (Anderson, 2018), and maps were produced using a combination of the *leaflet* (Cheng, Karambelkar, & Xie, 2017) and *tidycensus* (Walker, 2018) packages. Note that a github repository housing all the code and publicly available data from the project is available¹.

Effect size comparison. Student-level data from Oregon were used to empirically estimate achievement gap effect sizes by school. Across all effect size analyses, we pooled data across grades to estimate a single effect size, rather than estimating separate effects by grade. For the purposes of this investigation, we evaluated achievement gaps between students identifying as Hispanic/Latino and students identifying as White for all schools that met the minimum reporting size for both groups. Cohen's D was estimated as

$$d = \frac{\bar{X}_{foc} - \bar{X}_{ref}}{\sqrt{\frac{(n_{foc}-1)\sigma_{foc} + (n_{ref}-1)\sigma_{ref}}{n_{foc} + n_{ref} - 2}}} \quad (2)$$

where foc represents the focal group, in this case students identifying as Hispanic/Latino, and ref represents the reference group (students identifying as White). The numerator represents the difference in the means of the two distributions, while the denominator represents the pooled standard deviation. The V_c effect size was then estimated with the same data, and for the same set of schools. Comparisons between V_c and D were assessed both in terms of the correlation between the measures, as well as in terms of how discrepant V_c was from D .

Following the evaluation of V_c with D , we collapsed the data into counts by PLC within each school - i.e., we manually coarsened the data. We then estimated V_d and compared these estimates with both V_c and D . As stated previously, we expected V_c and D to differ marginally, given the different assumptions of the estimators. However, any differences in V_c

¹see https://github.com/DJAnderson07/ncme_18

and V_d could be interpreted as differences that arose due to the coarsening of the data.

Mapping school estimates. Publicly available data reporting on the percentage of students scoring in each PLC were obtained for California and Oregon, as described previously. Both files reported percentages by grade and across grades, by student subgroups. Similar to the student-level analyses, we used the pooled percentages across grades, rather than conducting separate analyses by grade. The Black-White achievement gap was estimated for California, while the Hispanic-White achievement gap was estimated for Oregon. In both cases, schools were only included if they reported percentages for both groups (i.e., met the minimal reporting threshold).

State level evaluations were followed up with county-level investigations using geographic and census data. Specifically, we obtained data on (a) the number of individuals identifying as Black (for California) or Hispanic/Latino (for Oregon), (b) the median housing cost, and (c) the number of individuals living at or below the poverty line within each county's census tracts, obtained from the 5-year American Community Survey (ACS) - data collected as part of the nationwide census (United States Census Bureau, 2016). Census tracts represent geographic areas with between approximately 1,200 to 8,000 people, with an optimum size of 4,000 people. Geographic boundaries of the census tracts were also obtained from the ACS survey. Maps were then produced with the census tracts, colored according to the respective variable, with schools displayed as "pins" on the map. Schools were colored according to the magnitude of the estimated achievement gap, and the relation between achievement gaps at the school level and demographic characteristics of the surrounding area were visually examined. All maps were interactive, helping to facilitate exploration (e.g., zoom, pan, identify schools by click). In our results, we share screen captures from these maps, although it is recognized that some of the features are difficult to portray through static images.

Results

Effect Size Comparisons

Figure 1 displays the relation between V_c and D both in terms of the bivariate relation and the distribution of differences with respect to V . The vertical line on the histogram indicates the mean of the differences. Note that it was slightly negative ($M = -0.12$), indicating that, on average for this sample, V represented a slightly lower estimate of the effect size.

References

- Anderson, D. (2018). *Esviz: Visualization and estimation of effect sizes*. Retrieved from <https://github.com/DJAnderson07/esviz>
- Bowen, N. K., & Bowen, G. L. (1999). Effects of crime and violence in neighborhoods and schools on the school behavior and performance of adolescents. *Journal of Adolescent Research, 14* (3), 319–342.
- California Assessment of Student Performance and Progress. (2017). *Test results for english language arts/literacy and mathematics*. Retrieved March 24, 2018, from <https://caaspp.cde.ca.gov/sb2017/ResearchFileList>
- California Department of Education. (2017). *Public schools and districts files*. Retrieved March 24, 2018, from <https://www.cde.ca.gov/ds/si/ds/pubschls.asp>
- Cheng, J., Karambelkar, B., & Xie, Y. (2017). *Leaflet: Create interactive web maps with the javascript 'leaflet' library*. Retrieved from <https://CRAN.R-project.org/package=leaflet>
- Dawes, S. S., Vidiyasa, L., & Parkhimovich, O. (2016). Planning and designing open government data programs: An ecosystem approach. *Government Information Quarterly, 33*(1), 15–27.
- Gonzales, N. A., Cauce, A. M., Friedman, R. J., & Mason, C. A. (1996). Family, peer, and neighborhood influences on academic achievement among african-american adolescents: One-year prospective effects. *American Journal of Community Psychology, 24* (3), 365.
- Google. (2018, March 15). *Google maps api*. Retrieved March 24, 2018, from <https://developers.google.com/maps/documentation/geocoding/start>
- Gregory, A., Skiba, R. J., & Noguera, P. A. (2010). The achievement gap and the discipline gap: Two sides of the same coin? *Educational Researcher, 39*(1), 59–68.
- Hanushek, E. A., Kain, J. F., & Rivkin, S. G. (2009). New evidence about brown v. Board of education: The complex effects of school racial composition on achievement. *Journal of Labor Economics, 27*(3), 349–383.
- Ho, A. D. (2008). The problem with “proficiency”: Limitations of statistics and policy under no

child left behind. *Educational Researcher*, 37(6), 351–360.

Ho, A. D. (2009). A nonparametric framework for comparing trends and gaps across tests.

Journal of Educational and Behavioral Statistics, 34(2), 201–228.

Ho, A. D., & Reardon, S. F. (2012). Estimating achievement gaps from test scores reported in ordinal “proficiency” categories. *Journal of Educational and Behavioral Statistics*, 37(4), 489–517.

Holland, P. W. (2002). Two measures of change in the gaps between the cdfs of test-score distributions. *Journal of Educational and Behavioral Statistics*, 27(1), 3–17.

Institute of Education Sciences. (2016, October). *IES policy regarding public access to research*.

Retrieved March 24, 2018, from <https://ies.ed.gov/funding/researchaccess.asp>

Kittrie, E., & Chattopadhyay, S. (2016). *NIH makes data sharing repositories publically viewable on healtdata.gov*. Retrieved March 22, 2018, from

https://datascience.nih.gov/Blog_HealthData.gov

Lee, M., & Madyun, N. (2009). The impact of neighborhood disadvantage on the black–white achievement gap. *Journal of Education for Students Placed at Risk*, 14(2), 148–169.

McCoy, D. C., Roy, A. L., & Sirkman, G. M. (2013). Neighborhood crime and school climate as predictors of elementary school academic quality: A cross-lagged panel analysis.

American Journal of Community Psychology, 52(1-2), 128–140.

National Science Foundation. (2016). *The national science foundation open government plan* (No. MSU-CSE-06-2).

No Child Left Behind. (2002). Act of 2001, 20 usca 6301 et seq.

Oregon Department of Education. (2017a). *Student assessment*. Retrieved March 24, 2018, from

<http://www.oregon.gov/ode/educator-resources/assessment/Pages/Assessment-Group-Reports.aspx>

Oregon Department of Education. (2017b). Retrieved March 24, 2018, from

<http://www.ode.state.or.us/pubs/labels/SchoolMail.xls>

R Core Team. (2017). *R: A language and environment for statistical computing*. Vienna,

Austria: R Foundation for Statistical Computing. Retrieved from

<https://www.R-project.org/>

Reardon, S. F., & Ho, A. D. (2015). Practical issues in estimating achievement gaps from coarsened data. *Journal of Educational and Behavioral Statistics*, 40(2), 158–189.

Sieber, R. E., & Johnson, P. A. (2015). Civic open data at a crossroads: Dominant models and current challenges. *Government Information Quarterly*, 32(3), 308–315.

Sirin, S. R. (2005). Socioeconomic status and academic achievement: A meta-analytic review of research. *Review of Educational Research*, 75(3), 417–453.

United States Census Bureau. (2016). *American community survey*. Retrieved March 24, 2018, from <https://www.census.gov/programs-surveys/acs/>

Walker, K. (2018). *Tidycensus: Load us census boundary and attribute data as 'tidyverse' and 'sf'-ready data frames*. Retrieved from <https://github.com/walkerke/tidycensus>

Wickham, H. (2017). *Tidyverse: Easily install and load the 'tidyverse'*. Retrieved from <https://CRAN.R-project.org/package=tidyverse>

Zuiderwijk, A., & Janssen, M. (2014). Open data policies, their implementation and impact: A framework for comparison. *Government Information Quarterly*, 31(1), 17–29.

Table 1

Content Area	ethniccd	n	mean	sd
Mathematics	Am. Indian	6341	217.81	21.67
Mathematics	Asian	13986	229.35	21.79
Mathematics	Black	9689	215.26	21.87
Mathematics	Hispanic	91967	217.60	19.61
Mathematics	Multiethnic	19646	222.63	20.39
Mathematics	Pac Islander	2760	218.01	19.99
Mathematics	White	234136	223.74	19.64
Reading/Language Arts	Am. Indian	6006	217.47	21.47
Reading/Language Arts	Asian	13941	225.01	19.27
Reading/Language Arts	Black	9237	216.00	21.07
Reading/Language Arts	Hispanic	89071	216.40	19.44
Reading/Language Arts	Multiethnic	18649	222.51	19.68
Reading/Language Arts	Pac Islander	2706	217.23	19.35
Reading/Language Arts	White	220974	223.48	19.26

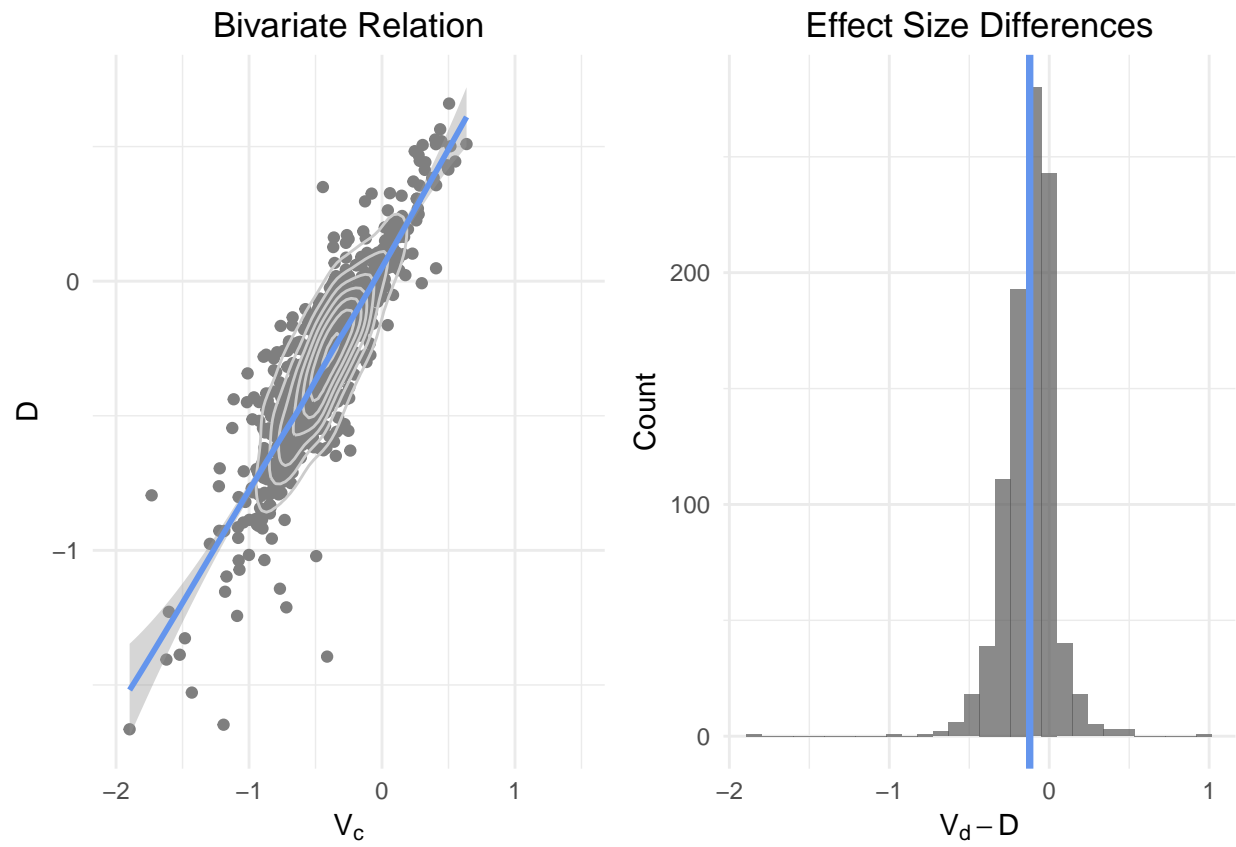


Figure 1. Differences in Cohen's D and V estimated with continuous data